

ECOMAss2

NathanAung

April 18, 2017

```
knitr::opts_chunk$set(echo = TRUE)
set.seed(42)
library(R6) #using R6 oop system
library(AER)
```

```
## Loading required package: car
## Loading required package: lmtest
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
## Loading required package: survival
```

```
library(stargazer)
```

```
##
## Please cite as:
## Hlavac, Marek (2015). stargazer: Well-Formatted Regression and Summary Statistics Tables.
## R package version 5.2. http://CRAN.R-project.org/package=stargazer
```

Q1

Define two classes that generates results:

```
#Class that generates all combinations in a vector
regMCLoop = R6Class(
  public = list(
    nrange=NULL,
    sigmarange=NULL,
    Urange=NULL,
    alrange=NULL,
    a2range=NULL,
    results = NULL,

    initialize = function(nrange,sigmarange,Urange,a2range=0,alrange=1)
    {
      self$nrangle = nrange
      self$sigmarange = sigmarange
      self$Urange = Urangle
      self$a2range = a2range
    }
  )
)
```

```

        self$a1range = a1range
        self$results = vector("list",length(nrange)*length(sigmarange)*length(Urange)*length(a2range))#list
    },

loop = function( a3=1, b1=1, b2=1, b3=0,IV=FALSE)
{
    i=1#elements of result vector
    for(n in self$nrange)
    {
        for(sigma in self$sigmarange)
        {
            for(dist in self$Urange)
            {
                for(a2 in self$a2range)
                {
                    for(a1 in self$a1range)
                    {
                        self$results[[i]] = regMClass$new(ns=n,sigma=sigma,dist=dist, a2=a2, a3=a3, b1=b1, b2=b2, b3=b3, IV=IV)
                        i=i+1#next result
                    }
                }
            }
        }
    }
}

)
)

```

```

regMClass = R6Class(
  public = list(
    bias = NULL,
    sd = NULL,
    CovRate=NULL,
    CovLength=NULL,
    bias2 =NULL,
    sd2 = NULL,
    CovRate2 = NULL,
    CovLength2 = NULL,
    ns = NULL,
    sigma= NULL,
    dist= NULL,
    a2 =NULL,
    a1 =NULL,
    fittedvalues = NULL,

initialize = function(ns = 20, n=1000, b1=1, b2=1, b3=0, a1=1, a2=0, a3=1, sigma=1,dist="norm", IV=FALSE)
{
    intervalrange =vector(mode="numeric",length=0)
    intervalrange2 =vector(mode="numeric",length=0)
    b1list = vector(mode="numeric",length=0)
    b2list = vector(mode="numeric",length=0)

```

```

inInterval=vector(mode="logical",length=0)
inInterval2=vector(mode="logical",length=0)#declare vectors so that append method works
for(j in 1:n)
{
  #__init__
  #we dont use a multivariate normal to generate it as this is equivalent.(independent and jointl
  Z1 = rnorm(n=ns,mean=0,sd=1)
  X2 = rnorm(n=ns,mean=0,sd=1)
  X3 = rnorm(n=ns,mean=0,sd=1)
  V = rnorm(n=ns,mean=0,sd=1)

  if (dist == "norm")
  {
    U = rnorm(n=ns,mean=0,sd=1)
  }
  else
  {
    U = rlnorm(n=ns,mean=0,sd=1)
    U=(U-mean(U))/sd(U)#z score rescaled
  }
  #Generate X1 and Y
  X1 = a1*Z1 + a2*X2 + a3*X3 +V
  X1 = X1/sd(X1) #renormalising so all values have sd of 1
  Y = b1*X1 + b2*X2 +b3*X3 + sigma*U
  #Estimates
  if (IV==FALSE){
    hat = lm(Y ~ X1+X2)
  }

  else{
    hat = ivreg(Y~X1+X2|X2+Z1)
  }
  b1list=append(b1list,hat$coefficients["X1"])
  b2list=append(b2list,hat$coefficients["X1"])
  #In confidence interval?
  #b1
  interval = confint(hat,parm = "X1",interval="confidence")
  inInterval = append(inInterval,b1<interval[2]&b1>interval[1])
  intervalrange = append(intervalrange,interval[2]-interval[1])
  #b2
  interval2 = confint(hat,parm = "X2",interval="confidence")
  inInterval2 = append(inInterval2,b2<interval2[2]&b2>interval[1])
  intervalrange2 = append(intervalrange2,interval2[2]-interval2[1])
  fittedvalues = hat$fitted.values

  #note that this could be probably be optimised greatly by a running mean,sd tally (saves memory

}
self$ns = ns # so we can sort by values used
self$sigma = sigma
self$dist = dist
self$a2 = a2
self$a1 = a1

```

```

        self$bias = mean(b1list)-1
        self$sd = sd(b1list)
        self$CovRate = mean(inInterval)
        self$CovLength = mean(intervalrange)

        self$bias2 = mean(b2list)-1
        self$sd2 = sd(b2list)
        self$CovRate2 = mean(inInterval2)
        self$CovLength2 = mean(intervalrange2)

        self$fittedvalues = fittedvalues
    }
)
)

```

We use a nested loop to generate all combinations needed and store them in a vector of results. These results are then stored within an object. It is possible that a recursive function approach could have been more elegant but for the number of combinations here, nested loop should suffice.

```

dataQ1 = regMCloop$new(nrange=c(20,200,2000),sigmarange=c(1,2),Urange=c("norm","lognorm"))
dataQ1$loop()

```

Here we create an instance of the class and use the loop methods to generate the results. This does take some time (~20s) as we are looping over quite a large number of iterations. We shall show the code here for this instance but as the code is largely similar, we will not show it for the sake of cluttering in following sections. This probably could have been a method for the regMCloop class.

```

Bias = vector(length=12)
SD = vector(length=12)
CovRate = vector(length=12)
CovLength = vector(length=12)
for(i in 1:12){
  Bias[i] = (dataQ1$results[[i]]$bias)
  SD[i] = (dataQ1$results[[i]]$sd)
  CovRate[i] = (dataQ1$results[[i]]$CovRate)
  CovLength[i]=(dataQ1$results[[i]]$CovLength)
}
ndf=c(rep(20,4),rep(200,4),rep(2000,4))
sigma = c(rep(1,2),rep(2,2),rep(1,2),rep(2,2),rep(1,2),rep(2,2))
Distribution = rep(c("Normal","LogNormal"),6)
Q1=data.frame(ndf,sigma,Distribution,Bias,SD,CovRate,CovLength)
stargazer(Q1,header=FALSE,type="latex",summary=FALSE)

```

a)

In all combinations, $|\text{bias}|$ is unanimously within ± 0.005 of 0. Considering that even under all combinations in which the number of samples is 20 that this is present, these results strongly indicate that the estimator is or very close to unbiased. It is impossible to say definitively due to the stochastic nature of the simulation, however the weak law of large numbers and central limit theorem point to this conclusion. This seems appropriate considering the theoretical aspects of the estimator. This is explained by the Gauss Markov Theorem, which guarantees that the OLS is the best unbiased linear estimator. In this case, all the conditions for the Gauss Markov Theorem are satisfied - the errors have mean 0, are homoskedastic (covariance matrix is identity) and are uncorrelated (generated independently) and so our estimator becomes the BLUE.

Table 1:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	20	1	Normal	0.010	0.242	0.949	0.987
2	20	1	LogNormal	-0.003	0.235	0.956	0.997
3	20	2	Normal	0.003	0.477	0.958	1.925
4	20	2	LogNormal	0.018	0.477	0.946	1.996
5	200	1	Normal	-0.005	0.068	0.957	0.280
6	200	1	LogNormal	0.003	0.071	0.954	0.280
7	200	2	Normal	-0.002	0.146	0.948	0.560
8	200	2	LogNormal	0.008	0.148	0.939	0.560
9	2,000	1	Normal	0.001	0.022	0.946	0.088
10	2,000	1	LogNormal	0.001	0.022	0.953	0.088
11	2,000	2	Normal	0.001	0.045	0.944	0.175
12	2,000	2	LogNormal	0.001	0.044	0.950	0.176

b)

For each combination of σ and probabilistic error distribution, we vary the number of samples and observe the results. In nearly all cases, bias decreases monotonically throughout. In the few cases where it does not, bias is very close to zero and is still within the same order of magnitude of the previous estimate. This is to be expected even with a consistent estimator as the estimator converges to the true value of the variable asymptotically. Perhaps with a higher number of simulations, this error could be reduced further. Testing across a larger number of sample sizes may have also graphically indicated the monotonicity of the bias as sample size increased. With only three data points, there could be any number of small changes in between.

c)

Coverage rate is consistently within 0.15 of the expected 95% coverage. Coverage rates between normal and log normal distributions appear to be inconclusive - that is there seems to be little difference. Similarly, σ appears to have limited effect on coverage rate. Higher standard deviation appears to slightly negatively affect coverage rate though the effect is again quite minute. The number of samples is also inconclusive - coverage rate appears to be uniform across n. One could of course perform an anova test to show these and reveal possible interaction effects.

d)

For all combinations, log normal would have generated much large confidence intervals than its normal counterpart.

This can be seen by examining the variance of normal and log normal distributions. With standard deviation σ , the variance for normal and lognormal respectively is:

$$\sigma^2$$

$$(e^{\sigma^2} - 1)(e^{2\mu + \sigma^2})$$

However because we z-score standardised the data in some cases it may be more accurate than the normal case - z-score scaling skews the data as every lognormal error has mean 0 and standard deviation 1. Obviously, if one draws from a distribution, not every sample will have this present. Initially, the interval length appears to be greater for lognormal distributions however in sample sizes greater than 20, the length is identical with their normal counterpart. This could be an indicator of slower convergence but again, the effect appears to be limited.

For unitary σ , the interval length is far smaller than their $\sigma = 2$ counterparts.

Sample size reduces interval length drastically for all combinations. This is because of the way confidence intervals are calculated - the standard error decreases with n and therefore so does the interval length.

2a)

```
dataQ2a = regMClloop$new(nrange=c(20,200,2000),sigmarange=c(1,2),Urange=c("norm","lognorm"))
dataQ2a$loop(b3=1,a3=1)
```

Table 2:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	20	1	Normal	0.555	0.323	0.594	1.279
2	20	1	LogNormal	0.583	0.322	0.579	1.281
3	20	2	Normal	0.549	0.514	0.813	2.129
4	20	2	LogNormal	0.576	0.515	0.834	2.150
5	200	1	Normal	0.572	0.099	0	0.361
6	200	1	LogNormal	0.576	0.094	0	0.361
7	200	2	Normal	0.580	0.158	0.035	0.604
8	200	2	LogNormal	0.574	0.153	0.035	0.606
9	2,000	1	Normal	0.577	0.029	0	0.113
10	2,000	1	LogNormal	0.579	0.030	0	0.113
11	2,000	2	Normal	0.578	0.049	0	0.190
12	2,000	2	LogNormal	0.578	0.049	0	0.189

It is clear that omitted variable bias is present here with an apparent consistent bias of 1. This would be ~0.33 if X1 was not rescaled.. We can calculate the theoretical bias here:

$$y = X_1\beta_1 + X_2\beta_2 + X_3\beta_3 + U$$

$$X_1 = \alpha_1 Z_1 + \alpha_3 X_3 + V_i$$

OLS estimates:

$$E(\hat{\beta}_1) = E \frac{\text{Cov}(X_1, y)}{\text{Var}(X_1)}$$

We use an approximation here as every X_1 value will be scaled slightly differently due to the method prescribed in rescaling the standard deviation to one.

$$E(\hat{\beta}_1) = E \frac{\text{Cov}(X_1, \frac{1}{\sqrt{3}}(X_1\beta_1 + X_2\beta_2 + X_3\beta_3 + U))}{1}$$

$$E(\hat{\beta}_1) = E \frac{\text{Cov}(X_1, X_1\beta_1) + \text{Cov}(X_1, X_2\beta_2) + \text{Cov}(X_1, X_3\beta_3) + \text{Cov}(X_1, U)}{\sqrt{3}}$$

$$E(\hat{\beta}_1) = E \frac{\beta_1 \text{Var}(X_1) + \beta_2 0 + \text{Cov}(Z_1 + X_3 + U, X_3\beta_3) + 0}{\sqrt{3}}$$

$$E(\hat{\beta}_1) = E(\beta_1 + \frac{\text{Cov}(Z_1 + X_3 + U, X_3\beta_3)}{\sqrt{3}})$$

$$E(\hat{\beta}_1) = E(\beta_1 + \frac{\beta_3\text{Cov}(Z_1, X_3) + \beta_3\text{Cov}(X_3, X_3) + \text{Cov}(U, X_3)}{\sqrt{3}})$$

$$E(\hat{\beta}_1) = \beta_1 + \frac{1}{\sqrt{3}}$$

$$\text{Bias} = \frac{1}{\sqrt{3}} \approx 0.5773$$

Standard deviation is also uniformly higher compared to a) and as such confidence intervals are larger. Coverage rate is especially poor here indicating a severe deviation from the requirements and assumptions of linear regression.

Whilst X_3 determines X_1 in both 1) and 2a), it does not determine Y in 1). As such, it does not fulfil the requirements for an omitted variable bias as endogeneity does not occur:

- 1) The omitted variable must be a determinant of the dependant variable (ie. regression coefficient is non-zero)
- 2) Omitted variable is correlated with the explanatory variable.

Whilst condition 2) is present through both questions, condition 1 only holds for question 2a). Because confidence length decreases with sample size and is centered around the mean of $\hat{\beta}$, the true beta is not even nearly within the range of the confidence interval for $\hat{\beta}$. Standard deviation is much larger here where the omitted variable bias is present. This is explained by the additional X_3 factor to y which appears as noise and adds variance within the data. This is more explicitly explained in the next question.

2b)

```
dataQ2b = regMClloop$new(nrange=c(20,200,2000),sigmarange=c(1,2),Urange=c("norm","lognorm"))
dataQ2b$loop(b3=1,a3=0)
```

Table 3:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	20	1	Normal	-0.012	0.337	0.954	1.392
2	20	1	LogNormal	-0.001	0.326	0.957	1.390
3	20	2	Normal	0.004	0.524	0.955	2.228
4	20	2	LogNormal	0.005	0.532	0.946	2.217
5	200	1	Normal	-0.004	0.102	0.949	0.396
6	200	1	LogNormal	-0.0003	0.099	0.955	0.396
7	200	2	Normal	0.001	0.165	0.943	0.625
8	200	2	LogNormal	-0.003	0.158	0.951	0.626
9	2,000	1	Normal	0.001	0.032	0.955	0.124
10	2,000	1	LogNormal	-0.0004	0.031	0.943	0.124
11	2,000	2	Normal	-0.001	0.049	0.956	0.196
12	2,000	2	LogNormal	-0.001	0.048	0.958	0.196

Bias will disappear as condition 1 is removed. There may be higher standard deviation than question one still as the variance of β can be calculated like so:

$$\text{Var}(\beta) = \text{Var}((X^T X)^{-1} X^T y)$$

$$\text{Var}(\hat{\beta}) = ((X^T X)^{-1} X^T) \sigma I ((X^T X)^{-1} X^T)^T$$

$$\text{Var}(\hat{\beta}) = (X^T X)^{-1} X^T \sigma X (X^T X)^{-1}$$

$$\text{Var}(\hat{\beta}) = \sigma (X^T X)^{-1}$$

And so depends on σ . In this case as X_3 is unobservable, it simply adds to the noise, thereby increasing variance.

Q3

a) Yes, it satisfies the two conditions required for an IV.

- 1) Z is not correlated with Y other than through X_1
- 2) Z is correlated relatively strongly with X_1

We know Z is not correlated with Y as it is generated independently.

b)

```
dataQ3 = regMClomp$new(nrange=c(20,200,2000),sigmarange=c(1,2),Urange=c("norm","lognorm"))
dataQ3$loop(b3=1,IV=TRUE)
```

Table 4:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	20	1	Normal	-0.265	2.868	0.949	11.274
2	20	1	LogNormal	-0.208	3.293	0.957	14.570
3	20	2	Normal	-0.190	2.635	0.961	10.513
4	20	2	LogNormal	-0.103	2.139	0.965	7.268
5	200	1	Normal	-0.003	0.180	0.945	0.694
6	200	1	LogNormal	-0.009	0.174	0.947	0.691
7	200	2	Normal	0.011	0.279	0.956	1.090
8	200	2	LogNormal	-0.013	0.281	0.949	1.099
9	2,000	1	Normal	-0.001	0.056	0.944	0.215
10	2,000	1	LogNormal	-0.0002	0.056	0.944	0.215
11	2,000	2	Normal	0.001	0.085	0.954	0.340
12	2,000	2	LogNormal	0.002	0.087	0.956	0.340

Asymptotic convergence is slower but bias is reduced eventually. IV estimates are notorious for having finite sample bias and thus for $n = 20$, bias is still very high. As n increases however the bias eventually appears to dissipate.

Standard deviations are higher due to higher presence of multicollinearity and use of IV estimates. IV estimates produce higher standard errors than ordinary OLS as it relies on more estimations. To demonstrate this:

$$\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{SST_x * R_{x,z}^2}$$

Under OLS though:

$$\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{SST_x}$$

Given $R_{x,z}^2$ is positive and between 0 and 1, the variance will always be smaller for OLS case holding other factors constant.

Interval lengths are also enlarged as a result of increased standard deviations.

4)

```
dataQ4 = regMClloop$new(nrange=200,sigmarange=1,Urange="norm",a1range=c(0.8,seq(0.6,0,by=-0.1)))
dataQ4$loop(b3=1,IV=TRUE)
```

Table 5:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	200	1	Normal	-0.013	0.213	0.954	0.824
2	200	1	Normal	-0.017	0.279	0.941	1.062
3	200	1	Normal	-0.023	0.329	0.954	1.297
4	200	1	Normal	-0.284	7.284	0.957	59.131
5	200	1	Normal	-0.168	1.951	0.951	8.340
6	200	1	Normal	0.070	10.446	0.951	206.402
7	200	1	Normal	-1.208	70.101	0.969	15,651.010
8	200	1	Normal	0.696	25.970	0.983	1,985.125

As the correlation between Z and X decreases, Z becomes takes on more and more of the characteristics of a weak instrument until it becomes an invalid instrument at 0. Bias increases accordingly, and the convergence of the estimator is slower. Because convergence for every given sample size increases with the strength of the IV, standard deviation at every sample size decreases accordingly. This has the follow over effect of increasing confidence interval lengths. Bias is also higher due to finite sample bias in IV regression.

Obviously when α_1 becomes zero it is longer an instrument and therefore the IVreg method is invalid. However when the instrumental variable is invalid, the bias and standard deviation curiously drops. This can be explained for a number of reasons which were outlined in the previous assignment. We again cite Morgan and Winship. Firstly, IV estimates can always be estimated as sample covariances are never exactly equal to zero (cannot draw an infinite sample). Therefore even for invalid instruments, Iv estimates can be computed. In fact, estimators for the standard errors are constructed under the assumption that the IV is valid and thereby generate “artificially small standard errors” (Morgan and Winship, 2015).

5)

```
dataQ5 = regMClloop$new(nrange=200,sigmarange=1,Urange="norm",a2range=seq(0,10,by=2))
dataQ5$loop()
```

Table 6:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	200	1	Normal	-0.0002	0.072	0.954	0.280
2	200	1	Normal	0.001	0.107	0.959	0.426
3	200	1	Normal	0.006	0.181	0.947	0.707
4	200	1	Normal	-0.006	0.265	0.938	1.010
5	200	1	Normal	0.011	0.335	0.953	1.321

β_1 does not appear to change with α_2 . At first this may seem surprising - common literature holds that multicollinearity increases standard errors.

However we see that this is only the case for β_2 :

The intuitive reasoning can be explained by thinking of the uncorrelated model($a_2=0$). In this case we have n points, in p dimensional space. ie. the column space of X. We seek to find the linear combination of these vectors which best explain y. A more precise way of defining this is we wish to find a project the data onto a hyperplane H such that the projection has smallest distance with y. In this case, it's a little simpler - we have a column space with rank 2.

Table 7:

	ndf	sigma	Distribution	Bias	SD	CovRate	CovLength
1	200	1	Normal	-0.0002	0.072	0.955	0.281
2	200	1	Normal	0.001	0.107	0.950	0.428
3	200	1	Normal	0.006	0.181	0.958	0.708
4	200	1	Normal	-0.006	0.265	0.961	1.013
5	200	1	Normal	0.011	0.335	0.968	1.326

In the correlated case, one of the columns is nearly linearly dependant on the other.

eg.

$$X_{1,new} = Z_1 + 5X_2 + U$$

In this case:

$$X_{1,new} \sim X_{1,old} + 4X_2$$

Then:

$$\hat{Y} = \beta_1 X_{1,new} + \beta_2 X_2$$

$$\hat{Y} = \beta_1 X_{1,old} + \beta_{1,old} 4X_2 + \beta_2 X_2$$

If we set

$$\beta_{2,new} = 1 - 4\beta_{1,old}$$

Then our estimate for \hat{Y} should be similar to our old one whilst keeping β_1 the same. Then β_1 would only vary in so much as X_1 would vary normally. However as β_2 has to compensate for the increasing value of α_2 , any variation in X_2 is magnified by the value of α_2 . Therefore as α_2 increases, β_2 continuously changes and the standard deviation increases correspondingly.

This reason why the old Y_{hat} is roughly the same value as the new one is because no new dimensions have really been added to the column space - X_2 is simply added again and as a projection is a linear combination, this can be compensated for very precisely by simply changing β_2 .

Therefore β_1 remains largely unchanged throughout changes in α_2 .